

empiricIST: An Integrated Software and analysis Tool for analyzing time-sampled sequence data such as EMPIRIC

README

empiricIST: An Integrated Software and analysis Tool for analyzing time-sampled sequence data such as EMPIRIC

Inês Fragata, Sebastian Matuszewski, Jeffrey D. Jensen and
Claudia Bank

September 4, 2017

Contents

| | | |
|----------|---------------------------------------|-----------|
| 1 | Introduction | 3 |
| 2 | Input data | 3 |
| 2.1 | empiricIST_MCMC_Input.py | 3 |
| 3 | Usage | 6 |
| 3.1 | Compilation | 6 |
| 3.2 | Execution | 9 |
| 4 | MCMC output | 11 |
| 4.1 | Combining Files | 16 |
| 4.2 | Concatenating Files | 18 |
| 4.3 | Visualization of trace data | 20 |
| 5 | DFE tail shape estimation | 20 |

1 Introduction

This program serves as an extension of the Bayesian Monte Carlo Markov chain (MCMC) method described in Bank et al. (2014) for estimating selection coefficients (growth rates) from engineered-mutation-driven experimental evolution data. These data are based on methods – such as EMPIRIC – in which specific mutations are engineered, introduced and compared against each other and a reference (e.g., wild-type) sequence. All mutants (and the wild-type) are assumed to have evolved together but independently in bulk competition over a number of generations with samples taken throughout the course of the experiment. Growth rate estimates are obtained from the number of reads obtained from deep sequencing. The motivation for the *empiricIST* software package is to provide an integrative software tool for the analysis of deep-mutational scanning data, and includes separate programs for processing raw sequence data, growth rate estimation, post-processing and analyses of growth rate estimates. Before using the software, please read the two accompanying papers by Bank et al. (2014) and ? that describe the methods and their underlying assumptions in more detail.

2 Input data

`empiricIST_MCMC` input data is expected to be in csv-format (i.e., comma-separated values) following a specific ordering and UNIX line endings. As part of the *empiricIST* software package, however, we provide a python script – `empiricIST_MCMC_Input.py` – that adjusts the raw data to match the specific input format needed. A minimal example of the raw data that is required to generate the `empiricIST_MCMC` input file is shown in Figure 1. Note that the raw data itself also needs to be csv-formatted with column entries separated by a comma (','), a semi-colon(';') or a tab ('\t'). Furthermore, the raw data file needs to have (exactly) one column either called 'sequence', 'Sequence', 'seq' or 'Seq' with at least two rows (the wild-type reference in the first row and a mutant) and at least additional 3 columns corresponding to the number of sequencing reads per sampled time point (and header cells giving the time of sampling).

2.1 `empiricIST_MCMC_Input.py`

The program `empiricIST_MCMC_Input.py` is written in Python (2.7) and serves as a link between the (raw) time-sampled sequence data (e.g., obtained from a deep-mutational scanning experiment) and the `empiricIST_MCMC` simulation program

| seq | 4.8 | 7.2 | 9.6 | 12 | 16.8 | 26.4 | 36 |
|-----------------------------------|-------|-------|-------|-------|-------|-------|--------|
| CCGGTCAAAACGGTTGGTCTGCTAACATGGAAA | 24901 | 28500 | 48710 | 58076 | 46121 | 52651 | 104330 |
| CCGGTAACAACGGTTGGTCTGCTAACATGGAAA | 626 | 738 | 1515 | 1497 | 1417 | 1928 | 2512 |
| CCGGTAAGAACGGTTGGTCTGCTAACATGGAAA | 579 | 499 | 1116 | 1510 | 1322 | 2080 | 3444 |
| CCGGTAATAACGGTTGGTCTGCTAACATGGAAA | 532 | 642 | 1198 | 1414 | 1151 | 1596 | 2210 |
| CCGGTACAACGGTTGGTCTGCTAACATGGAAA | 727 | 861 | 1721 | 1897 | 1752 | 2506 | 4040 |
| CCGGTACCAACGGTTGGTCTGCTAACATGGAAA | 1358 | 1536 | 2899 | 3046 | 3315 | 4906 | 7384 |
| CCGGTACGAACGGTTGGTCTGCTAACATGGAAA | 892 | 999 | 1979 | 2277 | 2194 | 3168 | 4896 |
| CCGGTACTAACGGTTGGTCTGCTAACATGGAAA | 880 | 1091 | 1957 | 2029 | 2112 | 3081 | 4727 |
| CCGGTAGAAACGGTTGGTCTGCTAACATGGAAA | 441 | 443 | 887 | 1235 | 1075 | 1645 | 2594 |
| CCGGTAGCAACGGTTGGTCTGCTAACATGGAAA | 505 | 633 | 1194 | 1631 | 1355 | 1948 | 2546 |
| CCGGTAGGAACGGTTGGTCTGCTAACATGGAAA | 418 | 431 | 907 | 1236 | 1082 | 1769 | 2924 |

Figure 1 – Schematic illustration of the minimal data needed to run the `empiricIST_MCMC` program.

for the estimation of mutant growth rates 'r'. While it primarily ensures that the input data matches the input format required by the MCMC simulation program, it comes with additional options that will be detailed here.

The general usage is as follows: After opening a command-line interface (e.g., Shell, Terminal) and navigating to the location of the `empiricIST_MCMC_Input.py` file, the program can be executed by typing

```
python empiricIST_MCMC_Input.py [options] .
```

Note that this requires that the 'PATHVARIABLE' for Python has been set correctly on your system. Please consult the online Python documentation for further details (<https://docs.python.org/2/>). Without specifying any options the program will exit with an error and provide a short documentation on its usage. To execute the program, the name of the (raw) data input file and the start of the sequencing read data needs to be specified (by invoking the '-f' and '-s' option, respectively). All options and their usage are given in Table 1.

Table 1 – A summary of the options of the `empiricIST_MCMC_Input.py` program.

| Short/Long option | Accepted values | Description |
|-------------------|-----------------|---|
| -h, -help | none | When the '-h'-option is invoked, a short documentation on the usage of the program is shown. Note that, if this option is invoked, the python program is not executed. |
| -f, -file= | string | The '-f'-option is a mandatory option, which passes the name of the (raw) data input file (csv formatted) to the python program. Files created by the python program will take the name of the input file and add option-dependent specific file identifiers. |
| -s, -skipcol= | integer | The '-s'-option is a mandatory option, which takes an integer value corresponding to the number of descriptive columns that precede the actual 'data matrix' of sequencing read counts. For example, for the raw data depicted in Figure 1, the user would have to pass '-s 1' (or equivalently '-skipcol=1'). Please note that the data matrix must always span all remaining columns. |

| | | |
|----------------|----------------------|---|
| | | When the '-o'-option is invoked, the python program will perform an outlier analysis. If '-o detect' (or equivalently '-outlier=detect'), the python program performs a log-linear regression analysis for all mutants. Data points are then classified as outliers on the basis of the DFBeta statistic with a cut-off value of 2 (with data points surpassing this cut-off regarded as outliers). For more details please consult Bank et al. (2014). If '-o impute' (or equivalently '-outlier=impute'), the python program performs a log-linear regression analysis for all mutants. Data points are then classified as outliers on the basis of the DFBeta statistic with a cut-off value of 2 and their studentized residuals with a cut-off value of 3 (with data points surpassing <i>both</i> cut-offs are regarded as outliers) and imputed as described in ?. Furthermore, an additional output file – whose name consists of 'ImputedData' along with the input file name – is produced that lists all imputed data points. In particular, the output is a simple matrix where rows denote different mutants and columns correspond to the different time points. An entry of '1' indicates that the data has <i>not</i> been imputed; an entry of '0' indicates that the data point has been imputed. |
| -o, -outlier= | 'detect' or 'impute' | |
| -l, -leadseq= | integer | When the '-l'-option is invoked, the first 'integer' characters, that precede the original mutant sequence (e.g., sites that function as DNA barcode or sequence tag), are removed. |
| -t, -trailseq= | integer | When the '-l'-option is invoked, the first 'integer' characters, that trail the original mutant sequence (e.g., sites that function as DNA barcode or sequence tag), are removed. |
| -p, -pool | none | When the '-p'-option is invoked, the DNA-sequences (characterizing the different mutants) are translated into amino acids. The data is then pooled based on their amino acid sequence, assuming that identical amino acid sequences, though differing in their DNA sequence (synonymous mutants), have the same growth rate (but different initial population sizes). Note that even if the '-p'-option is not invoked, data is pooled based on the sequence name (which can be any string and not only letters from the DNA alphabet). |
| -e, -exp | none | When the '-e'-option is invoked, time points are taken to be in hours instead of generations (i.e., the default). |
| -g, -group= | integer | When the '-g'-option is invoked, the data is grouped into subsets of mutants each of minimal size 'integer'. This results in more data sets with less mutants, such that the per-data set computation time is reduced, without affecting parameter estimates or the shape of the log-likelihood surface (compared to analysis of the full data set). The program also ensures, that mutants with identical mutant or protein ID (i.e., mutants that have an identical DNA- or amino acid sequence) remain in the same data sub-set as they are assumed to evolve under an identical growth rate (r). The name of the output file (i.e., the MCMC input file) is composed of the standard output file identifier 'MCMCInput', the grouping identifier (a consecutive number of the sub data sets created) and the name of the input file. |

| | | |
|--------------------------------------|--------------------------|---|
| <p><code>-i, --initialize</code></p> | <p><code>none</code></p> | <p>When the <code>'-i'</code>-option is invoked, and additional input file is created that specifies the initial growth rates (<code>r</code>) and initial population sizes (<code>c</code>) for all mutants based on the log-linear regression. Note that this could potentially bias the MCMC algorithm, since the initial starting point of the Markov chain could be trapped in a local log-likelihood optimum. Often, however, the median of the growth rates and the initial population sizes from the MCMC-DFE simulations are close enough to the corresponding estimates from the log-linear regression such that starting at these values could shorten the burn-in period and, thus, reduce the run time. For mutants with identical DNA or amino acid sequence, the mean initial population size is calculated (from these mutants) and taken as the starting value for the MCMC simulation program. Given the estimated mean initial population size, the growth rate is estimated from the log-linear regression. The name of the output file is composed of the standard output file identifier <code>'MCMCInput'</code>, an optional grouping identifier (see <code>'-g'</code>-option), the name of the input file, and an initialization file identifier <code>'_inputRC'</code>.</p> |
|--------------------------------------|--------------------------|---|

An illustration of the output file produced by the python program (i.e., the input file for the MCMC simulation program) is depicted in Figure 2.

Depending on the invoked options, the name of the MCMC simulation input file that is produced by the python program is given by the standard output file identifier `'MCMCInput'`, an optional grouping identifier (see `'-g'`-option) and the name of the raw data file.

3 Usage

We provide executables for Mac OS X, Windows and Linux. The C++ source code along with a system specific makefile are provided under a GNU General Public License as published by the Free Software Foundation. If you do not need to compile the program yourself you can skip the next subsection.

3.1 Compilation

Note that compilation requires that the Gnu Scientific Library (`gsl-library`) is installed on your system. Information on how to install the `gsl-library` can be found under <http://www.gnu.org/software/gsl/>. On Windows the easiest way to obtain the `gsl-library` is to install Cygwin (<http://www.cygwin.com/>) including the developers (all) packages. Alternatively, MinGW (<http://www.mingw.org/>) provides a "minimalist GNU for Windows" development environment. Under MinGW

| protID | seq | aa | r | rCIL | rCIU | s | sCIL | sCIU | 4.8 | 7.2 | 9.6 | 12 | 16.8 | 26.4 | 36x(4.8) | 4x(7.2) | 4x(9.6) | 4x(12) | 4x(16.8) | 4x(26.4) | 4x(36) |
|--------|-----------------------------|----------|--------------|--------------|--------------|--------------|---------------|--------------|-------|-------|-------|-------|-------|-------|----------|---------|---------|--------|----------|----------|--------|
| 1 | CAAAACGGTTGGTCTGCTAACATGAA | NGWSANME | 1 | 0.9959568244 | 1.0040431756 | 0 | 0.0040431756 | 0.0040431756 | 26082 | 29923 | 51373 | 61162 | 49083 | 56396 | 111318 | 1 | 1 | 1 | 1 | 1 | 1 |
| 2 | AACACGGTTGGTCTGCTAACATGAA | NGWSANME | 1.0226838373 | 1.0106965648 | 1.0346711097 | 0.0226838373 | 0.0106965648 | 0.0346711097 | 626 | 738 | 1515 | 1497 | 1417 | 1928 | 2512 | 1 | 1 | 1 | 1 | 1 | 0 |
| 3 | AATACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0001721595 | 1.0135174892 | 1.0206989296 | 0.0001721595 | 0.0135174892 | 0.0206989296 | 532 | 642 | 1196 | 1414 | 1151 | 1596 | 2210 | 1 | 1 | 1 | 1 | 1 | 0 |
| 4 | AAGACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0332578812 | 1.0241684213 | 1.042347341 | 0.0332578812 | 0.0241684213 | 0.042347341 | 579 | 499 | 1116 | 1510 | 1322 | 2080 | 3444 | 0 | 1 | 1 | 1 | 1 | 1 |
| 5 | AAAAACGGTTGGTCTGCTAACATGAA | NGWSANME | 1.039444547 | 1.0340523822 | 1.0448665173 | 0.039444547 | 0.0340523822 | 0.0448665173 | 717 | 706 | 1403 | 1759 | 1599 | 2344 | 4431 | 0 | 1 | 1 | 1 | 1 | 1 |
| 6 | AAACACGGTTGGTCTGCTAACATGAA | NGWSANME | 1.0311282594 | 1.0247358119 | 1.037494337 | 0.0311282594 | 0.0247358119 | 0.037494337 | 727 | 861 | 1721 | 1897 | 1752 | 2526 | 4040 | 1 | 1 | 1 | 1 | 1 | 0 |
| 7 | ACCAACGGTTGGTCTGCTAACATGAA | NGWSANME | 1.030305803 | 1.0188962326 | 1.0417553734 | 0.030305803 | 0.0188962326 | 0.0417553734 | 1358 | 1536 | 2899 | 3046 | 3315 | 4906 | 7384 | 1 | 1 | 1 | 1 | 1 | 1 |
| 8 | ACGACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.038054621 | 1.0293154247 | 1.0427838172 | 0.038054621 | 0.0293154247 | 0.0427838172 | 862 | 969 | 1979 | 2277 | 2194 | 3168 | 4896 | 1 | 1 | 1 | 1 | 1 | 1 |
| 9 | AACTACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0268965044 | 1.0174664707 | 1.0363352361 | 0.0268965044 | 0.0174664707 | 0.0363352361 | 880 | 1091 | 1957 | 2029 | 2112 | 3061 | 4727 | 1 | 1 | 1 | 1 | 1 | 1 |
| 10 | SAGAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0349594329 | 1.0214225183 | 1.0484963476 | 0.0349594329 | 0.0214225183 | 0.0484963476 | 441 | 443 | 887 | 1235 | 1075 | 1645 | 2594 | 1 | 1 | 1 | 1 | 1 | 1 |
| 11 | SAGAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0423237063 | 1.0293590295 | 1.0528638351 | 0.0423237063 | 0.0293590295 | 1.0528638351 | 418 | 451 | 907 | 1236 | 1082 | 1769 | 2924 | 1 | 1 | 1 | 1 | 1 | 1 |
| 12 | SGAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0314469396 | 1.0267708701 | 1.0371011172 | 0.0314469396 | 0.0267708701 | 0.0371011172 | 1285 | 1446 | 2717 | 3190 | 2893 | 4274 | 6744 | 1 | 1 | 1 | 1 | 1 | 1 |
| 13 | SGGACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0387027387 | 1.034055574 | 1.0433499033 | 0.0387027387 | 0.034055574 | 0.0433499033 | 1538 | 2438 | 4479 | 5834 | 5114 | 7841 | 14455 | 0 | 1 | 1 | 1 | 1 | 1 |
| 14 | SGGACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0452100346 | 1.0371389514 | 1.0533812179 | 0.0452100346 | 0.0371389514 | 1.0533812179 | 1210 | 1520 | 3143 | 3944 | 3525 | 5566 | 10297 | 1 | 1 | 1 | 1 | 1 | 1 |
| 15 | SGTACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0417924631 | 1.032683319 | 1.0509016071 | 0.0417924631 | 0.032683319 | 1.0509016071 | 1334 | 1680 | 3334 | 4422 | 3916 | 5776 | 10622 | 1 | 1 | 1 | 1 | 1 | 1 |
| 16 | SAGACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0382111179 | 1.0249112279 | 1.051510079 | 0.0382111179 | 0.0249112279 | 0.051510079 | 505 | 633 | 1194 | 1631 | 1355 | 1948 | 2546 | 1 | 1 | 1 | 1 | 1 | 0 |
| 17 | 6AGTACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0379829568 | 1.0279824283 | 1.0479824853 | 0.0379829568 | 0.0279824283 | 1.0479824853 | 367 | 504 | 907 | 1077 | 938 | 1493 | 1946 | 1 | 1 | 1 | 1 | 1 | 1 |
| 18 | 6TCAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0327954085 | 1.0248414625 | 1.0407493544 | 0.0327954085 | 0.0248414625 | 1.0407493544 | 620 | 624 | 1059 | 1258 | 1289 | 1843 | 3050 | 0 | 1 | 1 | 1 | 1 | 1 |
| 19 | 6TCAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0295706592 | 1.0114329184 | 1.0477084001 | 0.0295706592 | 0.0114329184 | 1.0477084001 | 923 | 979 | 1493 | 1770 | 1907 | 2960 | 4523 | 1 | 1 | 1 | 1 | 1 | 1 |
| 20 | 6TCAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0267647167 | 1.0151584922 | 1.0383710811 | 0.0267647167 | 0.0151584922 | 1.0383710811 | 746 | 815 | 1421 | 1550 | 1492 | 2385 | 3707 | 1 | 1 | 1 | 1 | 1 | 1 |
| 21 | 6TCTACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0275903062 | 1.0115202082 | 1.0436604043 | 0.0275903062 | 0.0115202082 | 1.0436604043 | 815 | 813 | 1418 | 1550 | 1741 | 2495 | 3859 | 1 | 1 | 1 | 1 | 1 | 1 |
| 22 | 7ATAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0402679414 | 1.0286828378 | 1.0518530451 | 0.0402679414 | 0.0286828378 | 1.0518530451 | 287 | 314 | 619 | 872 | 808 | 1123 | 2036 | 1 | 1 | 1 | 1 | 1 | 1 |
| 23 | 7ATACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.04127329 | 1.0308020441 | 1.051726156 | 0.04127329 | 0.0308020441 | 1.051726156 | 545 | 596 | 1111 | 1260 | 1305 | 2166 | 3645 | 1 | 1 | 1 | 1 | 1 | 1 |
| 24 | 7ATACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0519667794 | 1.0390486163 | 1.0648849425 | 0.0519667794 | 0.0390486163 | 1.0648849425 | 360 | 412 | 898 | 1140 | 1034 | 1637 | 2620 | 1 | 1 | 1 | 1 | 1 | 0 |
| 25 | 8ATACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0381992172 | 1.0284265966 | 1.0479718378 | 0.0381992172 | 0.0284265966 | 1.0479718378 | 441 | 504 | 994 | 1286 | 1133 | 1630 | 2620 | 1 | 1 | 1 | 1 | 1 | 0 |
| 26 | 8ATACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0175184847 | 1.0127712281 | 1.022267414 | 0.0175184847 | 0.0127712281 | 1.022267414 | 1867 | 2896 | 5528 | 5844 | 5477 | 7000 | 11436 | 0 | 1 | 1 | 1 | 1 | 1 |
| 27 | 9CATACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0204646084 | 0.0156041671 | 1.0253250497 | 0.0204646084 | 0.0156041671 | 1.0253250497 | 1423 | 1721 | 3295 | 3704 | 3254 | 4080 | 7371 | 1 | 1 | 1 | 1 | 1 | 1 |
| 28 | 10CCAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0322659712 | 1.0174185675 | 1.029113375 | 0.0322659712 | 0.0174185675 | 1.029113375 | 2334 | 3381 | 6081 | 6693 | 6273 | 8407 | 13853 | 1 | 1 | 1 | 1 | 1 | 1 |
| 29 | 10CCAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0257346803 | 1.0079968669 | 1.0434735036 | 0.0257346803 | 0.0079968669 | 1.0434735036 | 3548 | 5407 | 10407 | 9572 | 10159 | 14580 | 21039 | 1 | 1 | 1 | 1 | 1 | 1 |
| 30 | 10CCGACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0266821542 | 1.022197293 | 1.0311945791 | 0.0266821542 | 0.022197293 | 0.0311945791 | 2119 | 2453 | 4759 | 5527 | 5235 | 6761 | 11833 | 1 | 1 | 1 | 1 | 1 | 1 |
| 31 | 10CTACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0253921974 | 1.0172026128 | 1.033581782 | 0.0253921974 | 0.0172026128 | 1.033581782 | 2435 | 3193 | 6144 | 6538 | 6404 | 8626 | 13741 | 1 | 1 | 1 | 1 | 1 | 1 |
| 32 | 11CTAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0383530994 | 1.033589335 | 1.0591472652 | 0.0383530994 | 0.033589335 | 1.0591472652 | 1244 | 1528 | 2944 | 3465 | 3222 | 4516 | 8766 | 1 | 1 | 1 | 1 | 1 | 1 |
| 33 | 11CTAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0369712551 | 1.0308964095 | 1.0430461007 | 0.0369712551 | 0.0308964095 | 1.0430461007 | 1837 | 2246 | 4632 | 5323 | 4912 | 7328 | 12988 | 1 | 1 | 1 | 1 | 1 | 1 |
| 34 | 11CTGACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0374775488 | 1.0323016836 | 1.0426534141 | 0.0374775488 | 0.0323016836 | 1.0426534141 | 1111 | 1248 | 2399 | 3041 | 2805 | 4136 | 7475 | 1 | 1 | 1 | 1 | 1 | 1 |
| 35 | 11CTAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0375347989 | 1.0309096641 | 1.0436996336 | 0.0375347989 | 0.0309096641 | 1.0436996336 | 839 | 969 | 1905 | 2389 | 2224 | 2977 | 5973 | 1 | 1 | 1 | 1 | 1 | 1 |
| 36 | 11TTAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0337755821 | 1.0242591192 | 1.0432920449 | 0.0337755821 | 0.0242591192 | 1.0432920449 | 551 | 578 | 1064 | 1460 | 1332 | 1819 | 3330 | 1 | 1 | 1 | 1 | 1 | 1 |
| 37 | 11TTGACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0426843064 | 1.038460358 | 1.0498832548 | 0.0426843064 | 0.038460358 | 1.0498832548 | 522 | 624 | 1212 | 1636 | 1423 | 2192 | 4065 | 1 | 1 | 1 | 1 | 1 | 1 |
| 38 | 12GAAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0002616314 | 0.9840330148 | 1.0064602486 | 0.0002616314 | 0.0064602486 | 1.0064602486 | 625 | 660 | 1181 | 1424 | 1159 | 1355 | 1889 | 1 | 1 | 1 | 1 | 1 | 1 |
| 39 | 12GACACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0031157965 | 0.9952629826 | 1.0109686103 | 0.0031157965 | 0.00947370174 | 0.0109686103 | 501 | 522 | 965 | 1122 | 938 | 1085 | 1369 | 1 | 1 | 1 | 1 | 1 | 0 |
| 40 | 13GACACCGTTGGTCTGCTAACATGAA | NGWSANME | 0.951194181 | 0.9321370557 | 0.9714311804 | 0.0482158819 | 0.0678929443 | 0.026588196 | 557 | 745 | 1115 | 985 | 796 | 664 | 509 | 1 | 1 | 1 | 1 | 1 | 0 |
| 41 | 13GATACCGTTGGTCTGCTAACATGAA | NGWSANME | 0.959561097 | 0.9264346852 | 0.952695343 | 0.0604558903 | 0.0755653148 | 0.0473094657 | 402 | 483 | 636 | 696 | 474 | 376 | 252 | 1 | 1 | 1 | 1 | 1 | 0 |
| 42 | 14GCAACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.033424934 | 1.0222683662 | 1.0445815019 | 0.033424934 | 0.0222683662 | 1.0445815019 | 582 | 651 | 1124 | 1367 | 1405 | 2149 | 3302 | 1 | 1 | 1 | 1 | 1 | 1 |
| 43 | 14GCCACCGTTGGTCTGCTAACATGAA | NGWSANME | 1.0547949383 | 1.0487457244 | 1.0608441321 | 0.0547949383 | 0.0487457244 | 1.0608441321 | 1083 | 1474 | 2639 | 3364 | 3329 | 5378 | 7680 | 1 | 1 | 1 | 1 | 1 | 1 |

Figure 2 – Schematic illustration of the output produced by the python program (which serves as input for the MCMC simulation program). This input data has been created using the minimal raw data shown in Fig. 1, where the first five bases and the last base have been discarded (barcodes; '-l' and '-t' option). DNA sequences ('seq' column) have been translated to amino acids ('aa' column) and pooled ('-p' option), such that identical amino acid sequences have the same protein ID ('protID' column). Estimates of the growth rates 'r' ('r' column) and the selection coefficient 's' ('s' column) along with the 95%-confidence intervals are based on the log-linear regression (where 'rCIL' and 'rCIU' give the lower and upper boundary of the confidence interval for the growth rate r, respectively. Notation is analogous for the selection coefficient 's'). Please note, that while the MCMC simulation program assumes that mutants with identical sequence information (i.e., for sequences with identical 'protID') evolve at the same growth rate, log-linear estimates for 'r' and 's' are based on individual mutants. The columns '4.8', '7.2', '9.6', '12', '16.8', and '26.4' give the number of sequencing reads obtained from sampling at these time points for each mutant (row). If the '-o detect'-option is invoked, the matrix of sequencing reads is followed by an outlier matrix for the corresponding time points and mutants, where '0' indicate data points that were classified as outliers.

though, the gsl-library needs to be installed independently. A short instruction is given in the following paragraph.

Compilation of the program has successfully been tested on MacOSX (10.12.6) using 'clang' (version 6.0) and 'gcc' (version 4.9), on Ubuntu (14.04 LTS) using 'g++' (version 4.8.2), and on Windows (8.1) using 'g++' (version 4.8.1 under MinGW; version 4.9.2 under Cygwin).

Windows

Compilation of the program under Windows either requires Cygwin – i.e., a large collection of GNU and open source tools, which provide functionality similar to a Linux distribution on Windows – or MinGW – a complete open source programming tool set for the development of native Windows applications including both different compilers and a “minimal system” bourne shell command line interpreter system MSYS.

The easiest way to compile the program is by using the provided makefile. For that simply type

```
make
```

after having navigated to the folder where the source code is stored.

Please note that when using Cygwin commands and folder navigation are different to those from a native Windows Shell (e.g., Powershell) For example, under Cygwin you first need to type

```
cd /cygdrive/c
```

to navigate to the drive 'C:\'. Furthermore, note that Cygwin does not allow to link libraries statically.

When using MinGW/MSYS please check before compilation whether the gsl-library is installed. If not, a rough manual on how this can be done is given below. When already installed, the easiest way to compile the program is again to use the provided makefile. For that simply type

```
make
```

after having navigated to the folder where the source code is stored. Note that to navigate to the drive 'C:\' under MSYS one first has to type

```
cd /c .
```


Compilation and installation of the gsl-library from source Before compilation and installation of the gsl-library please read and follow the installation instructions provided with the most recent version of gsl.

1. Download the latest version of gsl from <http://ftpmirror.gnu.org/gsl/>
2. Navigate to the place where the downloaded tar archive is stored and unpack it by typing

```
tar -zxvf gsl-x.xx.tar.gz ,
```

where 'x.xx' should be replaced by the version number.

3. Navigate to 'gsl-x.xx/' and carefully read the 'INSTALL' document and follow the instructions to configure, make and install the gsl-library:

```
./configure  
make  
make install
```

4. GSL binaries, headers and library files are installed automatically in the 'bin/', 'include/gsl/', and 'lib/' subdirectories (if not specified otherwise; in that case you would also need to adjust the linker and compiler flags in the makefile).

Linux and MacOS X

The easiest way to compile the program is by using the provided makefile. For that simply use the shell and type

```
make
```

after having navigated to the folder where the source code is stored. Please note that you might want to adjust the makefile, in particular, to change the name of the executable which is by default set to '*empiricIST_MCMC.out*'.

3.2 Execution

Note that the *empiricIST_MCMC* program is a command line program which is run from a command-line interface (e.g., Shell, Terminal, Powershell), with arguments and parameters being passed over the command line. To ease execution of the program we provide OS specific scripts, where parameters and options can be specified by the user. An overview and description of the parameters and options

can be found in Table 2. Examples on how to manually execute the program are given when running the program without any parameters.

Under Windows there are two different scripts for running the program, depending on whether the program is run under Cygwin or using Powershell. Note that when using Cygwin, the source code needs to be compiled first (since Cygwin does not allow to compile statically) and 'PathToDataFile' needs to be preceded by '/cygdrive/c/' (assuming that the datafile is stored on the 'c' drive).

Table 2 – A summary of the parameters of the MCMC program.

| Category | Accepted values | Description |
|---------------|-------------------|---|
| Parameter | | |
| Data | | |
| -file | string | MANDATORY: Give the full path to the datafile (e.g., /users/me/PathToData/reads.csv). |
| -prefix | string | The name of the output file prefix (default DFE). The program will take the name of the input file and add the 'prefix', the time of execution and the file identifier (e.g., '_C'). This produces for example /PathToData/_ 'file' _ 'prefix' _ 'timeStamp' _C.txt. |
| -skipCol | integer, ≥ 0 | MANDATORY: Number of columns to skip in data file before read numbers start. |
| -outliers | | Pass the -outlier option if there is an outlier matrix in the data file. |
| MCMC | | |
| -burnin | integer, ≥ 0 | Number of accepted values that are discarded (burn-in period). During the burn-in period the parameters of the proposal distribution are optimized. By default 100,000 is used. |
| -subsampling | integer, ≥ 0 | After the burn-in period only every 'subsampling' accepted value is recorded (i.e., written to file). By default 1,000 is used. |
| -noSets | integer, ≥ 0 | Number of output data sets that are recorded each of size 'set'. By default 10 is used. |
| -set | integer, ≥ 0 | Number of recorded samples per set. By default 1,000 is used. |
| -growthRateSD | double, > 0 | Standard deviation of the proposal distribution of growth rates 'r' drawn from a Gaussian distribution. By default 0.0004 is used. |
| -popSizeSD | double, > 0 | Scale parameter of the proposal distribution of initial population sizes 'c' drawn from a Cauchy distribution. By default 0.0002 is used. |
| -initital | string | Specify the file (including its path) to the initializing data. An alternative way to initialize the growth rates 'r', the initial population sizes 'c' and to (optionally) set the parameters of the proposal distributions of the MCMC. If not specified growth rates are by default all set to 1 and initial population sizes correspond to the first observed read count. This option can for example be used to continue an MCMC run that has not been run long enough from the previous accepted sample. Note that in this case though, the burn-in has to be set to 0. |
| -hours | | If the -hours option is passed time points are assumed to be measured in hours. By default time is assumed to be measured in generations. |

| Output | | |
|----------------|-------------------|--|
| -logLTS | | If the -logLTS option is passed a time series of log-likelihoods is written to file. |
| -ESS | | If the -ESS option is passed the effective sample size (ESS) is calculated every 1000 accepted samples and written to file. |
| -screen | | If the -screen option is passed additional output will be written to screen. This option is mainly for inspection purposes. |
| Random numbers | | |
| -seed | integer, ≥ 0 | Sets the random number seed. By default the random number seed is created automatically based on computer run time. This option is mainly for inspection purposes. |

4 MCMC output

By default, the MCMC program outputs the raw growth rate and initial population sizes samples, along with two separate files containing summary and diagnostic statistics. The same output can optionally be generated for the log-likelihood data. Each file starts with a list of the input/parameters. For all files, data is written in tab-separated format, such that it be displayed nicely and easily with any spreadsheet application (such as Excel or OpenOffice).

Table 3 – A summary of the output of the MCMC program.

| File Parameter | Description |
|--------------------------|--|
| <u>.*_R</u> | |
| sample | Consecutive number of samples. Sample '0' gives the initial values. |
| r.* | Sampled value for the growth rate 'r' for all mutants. |
| <u>.*_C</u> | |
| sample | Consecutive number of samples. Sample '0' gives the initial values. |
| c.* | Sampled value for the initial population size 'c' for all mutants. |
| <u>.*_logLTS</u> | |
| sample | Consecutive number of samples. Sample '0' gives the initial values. |
| logL | Log-likelihood for the current sampled values for the initial population size 'c' and the growth rate 'r'. |
| <u>.*_R_quantiles</u> | |
| protID | Protein ID as specified by the input file. |
| mutant | Consecutive number of mutant identifier 'r.*'. |
| i% | Values for the $i\%$ -quantile of all samples of the growth rate 'r' for each mutant, where $i = 0, 1, 2.5, 5, 25, 50, 75, 95, 97.5, 99, 100$. |
| <u>.*_C_quantiles</u> | |
| protID | Protein ID as specified by the input file. |
| mutant | Consecutive number of mutant identifier 'c.*'. |
| i% | Values for the $i\%$ -quantile of all samples of the initial population size 'c' for each mutant, where $i = 0, 1, 2.5, 5, 25, 50, 75, 95, 97.5, 99, 100$. |
| <u>.*_logL_quantiles</u> | |
| logL | Log-Likelihood identifier. |
| i% | Values for the $i\%$ -quantile of all samples of the log-likelihood, where $i = 0, 1, 2.5, 5, 25, 50, 75, 95, 97.5, 99, 100$. |
| <u>.*_ess</u> | |
| sample | Number of samples after which effective sample size (ESS) is calculated. Note that the ESS is calculated every 'setSize' samples. |
| minESS | Minimum effective sample size computed for any parameter of interest (i.e., growth rate 'r', initial population size 'c' and log-likelihood). Note that $ESS \leq sample$. |
| r.* | ESS for growth rate 'r.*' for all mutant'. |
| c.* | ESS for initial population size 'c.*' for all mutants. |
| logL | ESS for log-likelihood. |
| acceptRatio | Overall acceptance ratio. To ensure high efficiency of the MCMC, the width of the proposal distributions should be chosen such that the acceptance ratio is between 0.15 – 0.45. Performance is maximal when the acceptance ratio is around 0.25. During the burn-in period the widths of the proposal distributions are automatically tuned to ensure that the acceptance is close to optimal, and efficiency is maximized during sampling. Thus, a sufficiently long burn-in period not only increases the chance that recorded samples are actually taken from the posterior distribution, but also that the width of the proposal distribution is set appropriately. |

| .*_Diag_R | |
|-----------|---|
| protID | Protein ID as specified by the input file. |
| mutant.* | Consecutive number of mutant identifier 'r.*'. |
| HD(*) | <p>Hellinger distance (HD) between sets of samples from two probability distributions. Note that HD is bounded by $0 \leq HD \leq 1$ and can be used to inspect the similarity between two distributions, where $HD = 0$ corresponds to no divergence and $HD = 1$ corresponds to no common support between the distributions. The HD can be used to diagnose the MCMC in terms of its burn-in and whether samples obtained at different points of time came (most likely) from the same (posterior) distribution. Note that one cannot determine if the MCMC chain has truly converged, but only if a chain is internally similar. Here, the HD is calculated for up to 10 equally sized sets of consecutive samples from the MCMC simulation. To obtain sufficient statistical power, the HD between two sets of samples is calculated only if each set consisted of at least 1000 samples. If the total number of samples is less than $10 \times 1000 = 10000$, the number of batches is chosen such that the total number of samples is divided into sets of samples of size 1000 each. If the total number of samples exceeds 10000, the number of samples per set is given by the total number of samples divided by 10 (i.e., the maximal number of batches). If the HD between sets of samples is less than 0.1 the distribution of posterior samples shows a high degree of similarity; if $0.1 \leq HD \leq 0.3$ the distribution of posterior samples are still quite similar, but may require closer inspection; if $0.3 \leq HD \leq 0.5$ sets of samples are vaguely similar and should be inspected more closely; a $HD > 0.5$ indicates strong dis-similarity between sets of samples and could be an indicator that all samples that were taken before might not be from the posterior distribution and should be discarded as burn-in. Note that the HD depends on the degree of autocorrelation between samples. Thus, a high HD might not necessarily indicate that samples were obtained from different sampling distributions, but poor mixing (i.e., a low ESS) for the parameter of interest. For details see Boone et al. (2014).</p> |
| mean | The mean of the posterior distribution for the parameter of interest. |
| SD | The standard deviation (SD) of the posterior distribution for the parameter of interest calculated with respect to the total number of samples. |
| median | The median of the posterior distribution for the parameter of interest. |
| 2.5% | The 2.5% quantile of the posterior distribution for the parameter of interest. |
| 97.5% | The 97.5% quantile of the posterior distribution for the parameter of interest. |
| ESS | The effective sample size for the parameter of interest. |
| minHD | The minimum HD calculated between consecutive batches. If there are not enough samples (more than 2000) to calculate the HD this field will read -1. |
| maxHD | The minimum HD calculated between consecutive batches. If there are not enough samples (more than 2000) to calculate the HD this field will read -1. |
| .*_Diag_C | |
| protID | Protein ID as specified by the input file. |
| mutant | Consecutive number of mutant identifier 'c.*'. |

| | |
|--------------|---|
| | <p>Hellinger distance (HD) between sets of samples from two probability distributions. Note that HD is bounded by $0 \leq HD \leq 1$ and can be used to inspect the similarity between two distributions, where $HD = 0$ corresponds to no divergence and $HD = 1$ corresponds to no common support between the distributions. The HD can be used to diagnose the MCMC in terms of its burn-in and whether samples obtained at different points of time came (most likely) from the same (posterior) distribution. Note that one cannot determine if the MCMC chain has truly converged, but only if a chain is internally similar. Here, the HD is calculated for up to 10 equally sized sets of consecutive samples from the MCMC simulation. To obtain sufficient statistical power, the HD between two sets of samples is calculated only if each set consisted of at least 1000 samples. If the total number of samples is less than $10 \times 1000 = 10000$, the number of batches is chosen such that the total number of samples is divided into sets of samples of size 1000 each. If the total number of samples exceeds 10000, the number of samples per set is given by the total number of samples divided by 10 (i.e., the maximal number of batches). If the HD between sets of samples is less than 0.1 the distribution of posterior samples shows a high degree of similarity; if $0.1 \leq HD \leq 0.3$ the distribution of posterior samples are still quite similar, but may require closer inspection; if $0.3 \leq HD \leq 0.5$ sets of samples are vaguely similar and should be inspected more closely; a $HD > 0.5$ indicates strong dis-similarity between sets of samples and could be an indicator that all samples that were taken before might not be from the posterior distribution and should be discarded as burn-in. Note that the HD depends on the degree of autocorrelation between samples. Thus, a high HD might not necessarily indicate that samples were obtained from different sampling distributions, but poor mixing (i.e., a low ESS) for the parameter of interest. For details see Boone et al. (2014).</p> |
| HD(*) | |
| mean | The mean of the posterior distribution for the parameter of interest. |
| SD | The standard deviation (SD) of the posterior distribution for the parameter of interest calculated with respect to the total number of samples. |
| median | The median of the posterior distribution for the parameter of interest. |
| 2.5% | The 2.5% quantile of the posterior distribution for the parameter of interest. |
| 97.5% | The 97.5% quantile of the posterior distribution for the parameter of interest. |
| ESS | The effective sample size for the parameter of interest. |
| minHD | The minimum HD calculated between consecutive batches. If there are not enough samples (more than 2000) to calculate the HD this field will read -1 . |
| maxHD | The minimum HD calculated between consecutive batches. If there are not enough samples (more than 2000) to calculate the HD this field will read -1 . |
| <hr/> | |
| .*_Diag_logL | |
| logL | Log-likelihood tag. |
| mutant | Consecutive number of mutant identifier 'c.*'. |

Hellinger distance (HD) between sets of samples from two probability distributions. Note that HD is bounded by $0 \leq HD \leq 1$ and can be used to inspect the similarity between two distributions, where $HD = 0$ corresponds to no divergence and $HD = 1$ corresponds to no common support between the distributions. The HD can be used to diagnose the MCMC in terms of its burn-in and whether samples obtained at different points of time came (most likely) from the same (posterior) distribution. Note that one cannot determine if the MCMC chain has truly converged, but only if a chain is internally similar. Here, the HD is calculated for up to 10 equally sized sets of consecutive samples from the MCMC simulation. To obtain sufficient statistical power, the HD between two sets of samples is calculated only if each set consisted of at least 1000 samples. If the total number of samples is less than $10 \times 1000 = 10000$, the number of batches is chosen such that the total number of samples is divided into sets of samples of size 1000 each. If the total number of samples exceeds 10000, the number of samples per set is given by the total number of samples divided by 10 (i.e., the maximal number of batches). If the HD between sets of samples is less than 0.1 the distribution of posterior samples shows a high degree of similarity; if $0.1 \leq HD \leq 0.3$ the distribution of posterior samples are still quite similar, but may require closer inspection; if $0.3 \leq HD \leq 0.5$ sets of samples are vaguely similar and should be inspected more closely; a $HD > 0.5$ indicates strong dis-similarity between sets of samples and could be an indicator that all samples that were taken before might not be from the posterior distribution and should be discarded as burn-in. Note that the HD depends on the degree of autocorrelation between samples. Thus, a high HD might not necessarily indicate that samples were obtained from different sampling distributions, but poor mixing (i.e., a low ESS) for the parameter of interest. For details see Boone et al. (2014).

HD(*)

mean The mean of the posterior distribution for the parameter of interest.

SD The standard deviation (SD) of the posterior distribution for the parameter of interest calculated with respect to the total number of samples.

median The median of the posterior distribution for the parameter of interest.

2.5% The 2.5% quantile of the posterior distribution for the parameter of interest.

97.5% The 97.5% quantile of the posterior distribution for the parameter of interest.

ESS The effective sample size for the parameter of interest.

minHD The minimum HD calculated between consecutive batches. If there are not enough samples (more than 2000) to calculate the HD this field will read -1 .

maxHD The minimum HD calculated between consecutive batches. If there are not enough samples (more than 2000) to calculate the HD this field will read -1 .

.*_Diag_summary

samples Absolute number of accepted samples taken during the MCMC run.

minESS(c) The minimum effective sample size (ESS) that was observed for any initial population size 'c.*'.

maxACT(c) The maximal auto-correlation time (ACT) that was observed for any initial population size 'c.*'.

minESS(r) The minimum effective sample size (ESS) that was observed for any growth rate 'r.*'.

maxACT(r) The maximal auto-correlation time (ACT) that was observed for any growth rate 'r.*'.

minESS(logL) The minimum effective sample size (ESS) that was observed for the log-likelihood.

maxACT(r) The maximal auto-correlation time (ACT) that was observed for the log-likelihood.

minESS(all) The minimum effective sample size (ESS) that was observed for all parameters.

maxACT(all) The maximal auto-correlation time (ACT) that was observed for all parameters.

| | |
|--------------|--|
| acceptRatio | The overall acceptance ratio of accepted (and recorded) samples. To ensure high efficiency of the MCMC, the width of the proposal distributions should be chosen such that the acceptance ratio is between 0.15 – 0.45. Performance is maximal with an acceptance ratio around 0.25. During the burn-in period the width of the proposal distributions is automatically tuned – based on the acceptance ratio – such that the acceptance ratio, when recording samples, has close to maximal efficiency. Thus, a sufficiently long burn-in period not only increases the chance that recorded samples are actually taken from the posterior distribution, but also that the width of the proposal distribution is set appropriately. |
| growthRateSD | Standard deviation of the proposal distribution of growth rates 'r' drawn from a Gaussian distribution after auto-tuning. |
| popSizeSD | Scale parameter of the proposal distribution of initial population sizes 'c' drawn from a Cauchy distribution after auto-tuning. |
| .*_initialRC | This file prints the last accepted values of the MCMC run so that these could be used as initial values, e.g., to continue an MCMC run that has not yielded enough independent samples. The first line gives the last sampled growth rates 'r.*', the second line gives the last sampled initial population sizes 'c.*', and the third line gives the standard deviation of the proposal distribution of growth rates 'r' (Normal distribution) and the scale parameter of the proposal distribution of initial population sizes 'c' (Cauchy distribution) after auto-tuning. |

4.1 Combining Files

In case the data set has been split into multiple subsets to enhance computational performance (see '-g' option in Tab. 1), we provide scripts to assemble the individual MCMC output files for each sub-data set into a single file (restricted to 'diagnostic', 'quantiles', and 'initial population size' and 'growth rate' files). `Combine_All.sh` is a wrapper which executes all 'Combine' scripts (i.e., `Combine_Diagnostic_C.sh`, `Combine_Diagnostic_R.sh`, `Combine_Quantiles_C.sh`, `Combine_Quantiles_R.sh`, `Combine_PopSizes_C.sh` and `Combine_GrowthRates_R.sh`). It furthermore removes the time stamp from all `empiricIST_MCMC` output files and (optionally) deletes the files that will not be combined. More information on all the script files is given in Table 4.

Please ensure that the scripts have the appropriate file permissions to perform the operations. File permissions can be adjusted by using the command-line and typing

```
chmod 755 <file> .
```


Table 4 – A summary of the 'combine files scripts'.

| Script | No. of passed variables | Description |
|-------------------------|-------------------------|--|
| Combine_All.sh | 6 | <p>This script removes the time stamp from all MCMC output files, executes all individual combine scripts, and (optionally) deletes the files that will not be combined. To execute the script type <code>./Combine_All.sh pathToPerlRename pathToData prefixFileName suffixFileName maxIndex deleteFiles</code>.</p> <p>pathToPerlRename: Provide full path to 'rename.pl'.</p> <p>pathToData: Provide full path to data folder.</p> <p>prefixFileName: Provide file name prefix.</p> <p>suffixFileName: Provide file name suffix.</p> <p>maxIndex: The number of sub-data sets the original data has been split-up to.</p> <p>deleteFiles: When set to '1' files that will not be combined will be deleted. Set to 0 otherwise.</p> <p>For an example see Figure 3.</p> <p>Information on the individual combine scripts can be found below.</p> |
| Combine_Diagnostic_C.sh | 4 | <p>This script combines all MCMC output files of type 'Diagnostic_C' of all sub-data sets into a single file. Note that the reference sequence will be deleted, since it does not contain any relevant information.</p> <p>To execute the script type <code>./Combine_Diagnostic_C.sh pathToData prefixFileName suffixFileName maxIndex</code>.</p> <p>pathToData: Provide full path to data folder.</p> <p>prefixFileName: Provide file name prefix.</p> <p>suffixFileName: Provide file name suffix.</p> <p>maxIndex: The number of sub-data sets the original data has been split-up to.</p> |
| Combine_Diagnostic_R.sh | 4 | <p>As above, but for growth rates 'R'.</p> |
| Combine_Quantiles_C.sh | 4 | <p>This script combines all MCMC output files of type 'Quantiles_C' of all sub-data sets into a single file. Note that the reference sequence will be deleted, since it does not contain any relevant information.</p> <p>To execute the script type <code>./Combine_Quantiles_C.sh pathToData prefixFileName suffixFileName maxIndex</code>.</p> <p>pathToData: Provide full path to data folder.</p> <p>prefixFileName: Provide file name prefix.</p> <p>suffixFileName: Provide file name suffix.</p> <p>maxIndex: The number of sub-data sets the original data has been split-up to.</p> |
| Combine_Quantiles_R.sh | 4 | <p>As above, but for growth rates 'R'.</p> |

| | | |
|--------------------------|---|---|
| Combine_PopSizes_C.sh | 4 | <p>This script combines all MCMC output files of type '_C' of all sub-data sets into a single file. Note that the reference sequence will be deleted, since it does not contain any relevant information. To execute the script type</p> <pre>./Combine_PopSizes_C.sh pathToData prefixFileName suffixFileName maxIndex .</pre> <p>pathToData: Provide full path to data folder. prefixFileName: Provide file name prefix. suffixFileName: Provide file name suffix. maxIndex: The number of sub-data sets the original data has been split-up to.</p> |
| Combine_GrowthRates_R.sh | 4 | As above, but for growth rates 'R'. |
| RemoveMCMCTimeStamp.sh | 3 | <p>This script removes the time stamp from all <code>empiricIST_MCMC</code> output files. To execute the script type</p> <pre>./RemoveMCMCTimeStamp.sh pathToPerlRename pathToData prefixFileName .</pre> <p>pathToPerlRename: Provide full path to 'rename.pl'. pathToData: Provide full path to data folder. prefixFileName: Provide file name prefix.</p> |

4.2 Concatenating Files

When an MCMC analysis has been continued across multiple runs (i.e., by starting another MCMC run from the last accepted sample; see `-initial` option above), the 'ConcatenateData.sh' script can be used to produce a single file of all sampled growth rates 'r', initial population sizes 'c' and log-likelihoods, respectively.

$$\underbrace{\text{MCMCInput}}_{\text{prefixName}} - \overbrace{99}^{\text{maxIndex}} - \underbrace{\text{HGW_VIE}}_{\text{suffixName}} \overbrace{\text{_08_06_2015_142000_}}^{\text{date stamp}} \underbrace{\text{R}}_{\text{type}} .txt$$

Figure 3 – Illustration of how to set the command line arguments for the combine scripts.

Table 5 – A summary of the 'combine files scripts'.

| Script | No. of passed variables | Description |
|--------------------|-------------------------|---|
| ConcatenateData.sh | 8 | <p>This script removes the time stamp from all MCMC output files, (optionally) deletes the files that will not be combined, and concatenates the individual files. To execute the script type</p> <pre>./ConcatenateData.sh pathToData pathToMove pathToPerlRename prefixFileName suffixFileName simIdentifier maxIndex deleteFiles .</pre> <p>pathToData: Provide full path to data folder. pathToMove: Provide full path where to move data. pathToPerlRename: Provide full path to 'rename.pl'. prefixFileName: Provide file name prefix. suffixFileName: Provide file name suffix. simIdentifier: Provide indicator for which data to concatenate: R: growth rates, C: initial population sizes, logLTS: log-likelihood maxIndex: The number of MCMC runs the original data has been split-up to. For an example see Figure 3. deleteFiles: When set to '1' files that will not be combined will be deleted. Set to 0 otherwise.</p> |

Once all samples have been combined into a single file, statistics and diagnostics can be calculated by using the provided python programs `empiricIST_MCMC_Statistics.py` and `empiricIST_MCMC_DiagSummary.py`, respectively. Taking the raw samples as input, `empiricIST_MCMC_Statistics.py` creates the `quantiles-` and the `Diag` files as if obtained from a single MCMC run (for detail see above). The python program `empiricIST_MCMC_DiagSummary.py` can then be used to create the `Diag_summary` file over all individual diagnostic statistics (see above for details). Note though, that unlike the file produced from a single MCMC run, the one created by the python program does not report the standard deviation of the proposal distribution of growth rates 'r' and initial population sizes 'c', nor the acceptance ratio. Furthermore, it can only be used when the individual log-likelihood, growth rate and initial population size files exist. Table 6 and table 7 give an overview of the options for each program.

Table 6 – A summary of the options of the `empiricIST_MCMC_Statistics.py` program.

| Short/Long option | Accepted values | Description |
|-------------------|-----------------|--|
| -h, -help | none | When the '-h'-option is invoked, a short documentation on the usage of the program is shown. Note that, if this option is invoked, the python program is not executed. |
| -f, -file= | string | The '-f'-option is a mandatory option, which passes the name of the data input file to the python program. Files created by the python program will take the name of the input file and insert <code>_Quantiles.txt</code> and <code>_Diag_</code> , respectively. |

Table 7 – A summary of the options of the `empiricIST_MCMC_DiagSummary.py` program.

| Short/Long option | Accepted values | Description |
|-------------------|-----------------|--|
| -h, -help | none | When the '-h'-option is invoked, a short documentation on the usage of the program is shown. Note that, if this option is invoked, the python program is not executed. |
| -f, -file= | string | The '-f'-option is a mandatory option, which passes the base name of the data input files to the python program. For example, if the file names are <code>DFE_Diag_logTS.txt</code> , <code>DFE_Diag_C.txt</code> and <code>DFE_Diag_R.txt</code> the base name to be passed is <code>DFE_Diag</code> . The output file created by the python program will take the name of the input file and prepend <code>_summary.txt</code> . |
| -s, -samples= | integer | The '-s'-option is a mandatory option, which passes the number of (posterior) samples to the python program. |

4.3 Visualization of trace data

Additionally, there is a shell script '`FormatTracer.sh`' that formats the posterior sample output file (e.g., containing the initial population size 'c' or the growth rate 'r') such that it can be read and analyzed by *Tracer* (Rambaut et al. 2014). *Tracer* is a graphical tool for visualization and diagnostics of MCMC output that, for instance, displays the posterior distribution and its credibility interval, calculates the effective sample size (ESS; note that values might be slightly different from those calculated by the `empiricIST_MCMC` program since we use a more accurate but computational more intensive algorithm), and shows the trace of the posterior samples. Note that the input file is expected to be formatted as the output file created from the `Combine_PopSizes_C.sh/Combine_GrowthRates_R.sh` script. Otherwise the provided `Create_TailShapeFileR.sh` script can be used to obtain a correctly formatted input file (please see below for details).

`./FormatTracer.sh pathToData fileName .`

pathToData: Provide full path to data folder.

fileName: Provide file name without file extension (e.g., '.txt').

5 DFE tail shape estimation

The growth rate posterior samples obtained from the `empiricIST_MCMC` program can be used to estimate the shape of the beneficial tail of the distribution of fitness effects (DFE). As part of the *empiricIST* software package, we provide a python

script – `empiricIST_MCMC_TailShape.py` that fits a generalized pareto distribution to the observed beneficial mutations by maximizing the log-likelihood with respect to the shape and scale parameter κ and ψ , respectively. Based on the shape parameter κ one can discriminate between three different domains of attraction – the Weibull, Gumbel and Fréchet domain – each corresponding to a different extreme value distribution. In biological terms, these different domains quantify the level of adaptedness of the organism in its (experimental) environment. In particular, the Gumbel domain ($\kappa = 0$; null model corresponding to 'normal' level of adaptedness) is characterized by an exponential tail, whereas the Weibull domain ($\kappa < 0$; better adapted) has lighter than exponential tails, and the Fréchet domain ($\kappa > 0$; less well adapted) has heavier than exponential tails. For more details please consult Beisel et al. (2007).

While the python script primarily estimates the DFE tail shape parameter κ , there are some additional options that will be detailed here.

The general usage is as follows: The program can be executed by typing

```
python empiricIST_MCMC_TailShape.py [options] .
```

Without specifying any options the program will exit with an error and provide a short documentation on its usage, as it requires the name of the data input file (by invoking the '-f' option). Note that the input file needs to be formatted in a specific way. If the data set had been split into multiple subsets (see '-g' option in Tab. 1), and has been re-assembled by the using the provided shell scripts (see 'Combining Files'), the input file is already correctly formatted and there is nothing that needs to be done. However, if the data set has been analyzed as a whole (i.e., without being split into multiple subsets) the file containing the posterior growth rate samples needs to be reformatted to match the required input format. This can be done by using the provided shell script `Create_TailShapeFileR.sh`, which creates the input file for the `empiricIST_MCMC_TailShape.py` program and can be executed by typing

```
./Create_TailShapeFileR.sh pathToData fileName .
```

pathToData: Provide full path to data folder.

fileName: Provide file name without file extension (e.g., '.txt').

Please ensure that the script has the appropriate file permissions to perform the operations. File permissions can be adjusted by using the command-line and typing

```
chmod 755 Create_TailShapeFile_R.sh .
```

All options and their usage of the `empiricIST_MCMC_TailShape.py` program are given in Table 8.

Table 8 – A summary of the options of the `empiricIST_MCMC_TailShape.py` program.

| Short/Long option | Accepted values | Description |
|-------------------|-----------------|---|
| -h, -help | none | When the '-h'-option is invoked, a short documentation on the usage of the program is shown. Note that, if this option is invoked, the python program is not executed. |
| -f, -file= | string | The '-f'-option is a mandatory option, which passes the name of the data input file to the python program. Files created by the python program will take the name of the input file and add option-dependent specific file identifiers. Note that even if a random data set is created (see 'r'-option), a file name must be provided since it serves as the prefix for the output file name. |
| -m, -missing | none | When the '-m'-option is invoked, the distribution of measured fitnesses is shifted relative to the smallest observed selection coefficient to account for missing data (i.e., selection coefficients too small to have been observed, though this might not be a problem with EMPIRIC data). |
| -s, -samples= | integer | When the '-s'-option is invoked, the python program will only consider samples with more than 'samples' beneficial mutations for maximum likelihood estimation. Note that when this option is not specified, the default is set to 10. |
| -l, -lhoodrt | none | When the '-l'-option is invoked, a likelihood-ratio test with null hypotheses $\kappa = 0$ against median($\hat{\kappa}$) is performed. Note that the distribution of the test statistic is generated by using a parametric bootstrap approach with 10.000 samples (Beisel et al. 2007, see). |
| -r, -random= | double | When the '-r'-option is invoked, 100 random data sets are created with 100 samples each drawn from a generalized pareto distribution with scale parameter $\psi = 1$ and κ passed as command line argument. |

By default the program will always create two output files that contain the κ and ψ estimates called `<InputFileName>_TailShape_KappaShape.txt` and `<InputFileName>_TailShape_PsiShape.txt`, respectively. Note that even if a random data set is created (see 'r'-option), a file name must be provided since it serves as the prefix for the output file name. When performing a likelihood-ratio test (by invoking the 'l'-option) an additional file called `<InputFileName>_LRT.txt` is created containing the $\hat{\kappa}$ against which $H_0 : \kappa = 0$ is evaluated. Furthermore, the output gives the maximum-likelihood estimate $\hat{\psi}_{\kappa_0}$, i.e., the estimated scale parameter ψ restricting $\kappa = 0$, the value of the test statistic

$$-2\log(\Lambda) = 2(\mathcal{L}(\mathbf{X} \mid \hat{\kappa}, \hat{\psi}) - \mathcal{L}(\mathbf{X} \mid 0, \hat{\psi}_{\kappa_0})), \quad (1)$$

where \mathbf{X} denotes a single vector of posterior samples of growth rates and \mathcal{L} is the log-likelihood function, and the associated p-value along with the sample size.

Note that power critically depends on sample size as has been discussed in Beisel et al. (2007) and del Castillo and Serra (2015) (the latter rather in the context of estimating κ accurately).

References

- Bank, C., R. T. Hietpas, A. Wong, D. N. Bolon, and J. D. Jensen, 2014. A bayesian mcmc approach to assess the complete distribution of fitness effects of new mutations: Uncovering the potential for adaptive walks in challenging environments. *Genetics* 196:841–852.
- Beisel, C. J., D. R. Rokyta, H. A. Wichman, and P. Joyce, 2007. Testing the extreme value domain of attraction for distributions of beneficial fitness effects. *Genetics* 176:2441–2449.
- Boone, E. L., J. R. Merrick, and M. J. Krachey, 2014. A hellinger distance approach to mcmc diagnostics. *Journal of Statistical Computation and Simulation* 84:833–849.
- del Castillo, J. and I. Serra, 2015. Likelihood inference for generalized pareto distribution. *Computational Statistics And Data Analysis* 83:116 – 128.
- Rambaut, A., M. Suchard, D. Xie, and A. Drummond, 2014. Tracer v1.6.